

# Data Manipulation Challenge

## A Mental Model for Method Chaining in Pandas

### Data Manipulation Challenge - A Mental Model for Method Chaining in Pandas

! Challenge Requirements In Section [Student Analysis Section](#)

- Complete all discussion questions for the seven mental models (plus some extra requirements for higher grades)

! Note on Python Usage

**Recommended Workflow: Use Your Existing Virtual Environment** If you completed the Tech Setup Challenge Part 2, you already have a virtual environment set up! Here's how to use it for this new challenge:

1. **Clone this new challenge repository** (see Getting Started section below)
2. **Open the cloned repository in Cursor**
3. **Set this project to use your existing Python interpreter:**
  - Press **Ctrl+Shift+P** → “Python: Select Interpreter”
  - Navigate to and choose the interpreter from your existing virtual environment (e.g., `your-previous-project/venv/Scripts/python.exe`)
4. **Activate the environment in your terminal:**
  - Open terminal in Cursor (‘Ctrl + ‘)
  - Navigate to your previous project folder where you have the `venv` folder
  - **Pro tip:** You can quickly navigate by typing `cd` followed by dragging the folder from your file explorer into the terminal
  - Activate using the appropriate command for your system:
    - **Windows Command Prompt:** `venv\Scripts\activate`

- **Windows PowerShell:** `.\venv\Scripts\Activate.ps1`
- **Mac/Linux:** `source venv/bin/activate`
- You should see `(venv)` at the beginning of your terminal prompt

5. **Install additional packages if needed:** `pip install pandas numpy matplotlib seaborn`

#### **Cloud Storage Warning**

**Avoid using Google Drive, OneDrive, or other cloud storage for Python projects!** These services can cause issues with: - Package installations failing due to file locking - Virtual environment corruption - Slow performance during pip operations

**Best practice:** Keep your Python projects in a local folder like `C:\Users\YourName\Documents\` or `~/Documents/` instead of cloud-synced folders.

**Alternative: Create a New Virtual Environment** If you prefer a fresh environment, follow the Quarto documentation: <https://quarto.org/docs/projects/virtual-environments.html>. Be sure to follow the instructions to activate the environment, set it up as your default Python interpreter for the project, and install the necessary packages (e.g. pandas) for this challenge. For installing the packages, you can use the `pip install -r requirements.txt` command since you already have the `requirements.txt` file in your project. Some steps do take a bit of time, so be patient.

**Why This Works:** Virtual environments are portable - you can use the same environment across multiple projects, and Cursor automatically activates it when you select the interpreter!

## **The Problem: Mastering Data Manipulation Through Method Chaining**

**Core Question:** How can we efficiently manipulate datasets using `pandas` method chaining to answer complex business questions?

**The Challenge:** Real-world data analysis requires combining multiple data manipulation techniques in sequence. Rather than creating intermediate variables at each step, method chaining allows us to write clean, readable code that flows logically from one operation to the next.

**Our Approach:** We'll work with ZappTech's shipment data to answer critical business questions about service levels and cross-category orders, using the seven mental models of data manipulation through `pandas` method chaining.

### AI Partnership Required

This challenge pushes boundaries intentionally. You'll tackle problems that normally require weeks of study, but with Cursor AI as your partner (and your brain keeping it honest), you can accomplish more than you thought possible.

**The new reality:** The four stages of competence are Ignorance → Awareness → Learning → Mastery. AI lets us produce Mastery-level work while operating primarily in the Awareness stage. I focus on awareness training, you leverage AI for execution, and together we create outputs that used to require years of dedicated study.

## The Seven Mental Models of Data Manipulation

The seven most important ways we manipulate datasets are:

1. **Assign:** Add new variables with calculations and transformations
2. **Subset:** Filter data based on conditions or select specific columns
3. **Drop:** Remove unwanted variables or observations
4. **Sort:** Arrange data by values or indices
5. **Aggregate:** Summarize data using functions like mean, sum, count
6. **Merge:** Combine information from multiple datasets
7. **Split-Apply-Combine:** Group data and apply functions within groups

## Data and Business Context

We analyze ZappTech's shipment data, which contains information about product deliveries across multiple categories. This dataset is ideal for our analysis because:

- **Real Business Questions:** CEO wants to understand service levels and cross-category shopping patterns
- **Multiple Data Sources:** Requires merging shipment data with product category information
- **Complex Relationships:** Service levels may vary by product category, and customers may order across categories
- **Method Chaining Practice:** Perfect for demonstrating all seven mental models in sequence

## Data Loading and Initial Exploration

Let's start by loading the ZappTech shipment data and understanding what we're working with.

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime, timedelta

# Load the shipment data
shipments_df = pd.read_csv(
    "https://raw.githubusercontent.com/flyaflya/persuasive/main/shipments.csv",
    parse_dates=['plannedShipDate', 'actualShipDate']
)

# Load product line data
product_line_df = pd.read_csv(
    "https://raw.githubusercontent.com/flyaflya/persuasive/main/productLine.csv"
)

# Reduce dataset size for faster processing (4,000 rows instead of 96,805 rows)
shipments_df = shipments_df.head(4000)

print("Shipments data shape:", shipments_df.shape)
print("\nShipments data columns:", shipments_df.columns.tolist())
print("\nFirst few rows of shipments data:")
print(shipments_df.head(10))

print("\n" + "="*50)
print("Product line data shape:", product_line_df.shape)
print("\nProduct line data columns:", product_line_df.columns.tolist())
print("\nFirst few rows of product line data:")
print(product_line_df.head(10))

```

Shipments data shape: (4000, 5)

Shipments data columns: ['shipID', 'plannedShipDate', 'actualShipDate', 'partID', 'quantity']

First few rows of shipments data:

	shipID	plannedShipDate	actualShipDate	partID	quantity
0	10001	2013-11-06	2013-10-04	part92b16c5	6
1	10002	2013-10-15	2013-10-04	part66983b	2
2	10003	2013-10-25	2013-10-07	part8e36f25	1
3	10004	2013-10-14	2013-10-08	part30f5de0	1
4	10005	2013-10-14	2013-10-08	part9d64d35	6

5	10006	2013-10-14	2013-10-08	part6cd6167	15
6	10007	2013-10-14	2013-10-08	parta4d5fd1	2
7	10008	2013-10-14	2013-10-08	part08cadf5	1
8	10009	2013-10-14	2013-10-08	part5cc4989	10
9	10010	2013-10-14	2013-10-08	part912ae4c	1

=====

Product line data shape: (11997, 3)

Product line data columns: ['partID', 'productLine', 'prodCategory']

First few rows of product line data:

	partID	productLine	prodCategory
0	part00005ba	line4c	Liquids
1	part000b57d	line61	Machines
2	part00123bf	linec1	Marketable
3	part0021fc9	line61	Machines
4	part0027e86	line2f	Machines
5	part002ed95	line4c	Liquids
6	part0030856	lineb8	Machines
7	part0033dfd	line49	Liquids
8	part0037a2a	linea3	Marketable
9	part003caee	linea3	Marketable

### Understanding the Data

**Shipments Data:** Contains individual line items for each shipment, including: - **shipID:** Unique identifier for each shipment - **partID:** Product identifier - **plannedShipDate:** When the shipment was supposed to go out - **actualShipDate:** When it actually shipped - **quantity:** How many units were shipped

**Product Category and Line Data:** Contains product category information: - **partID:** Links to shipments data - **productLine:** The category each product belongs to - **prodCategory:** The category each product belongs to

**Business Questions We'll Answer:** 1. Does service level (on-time shipments) vary across product categories? 2. How often do orders include products from more than one category?

## The Seven Mental Models: A Progressive Learning Journey

Now we'll work through each of the seven mental models using method chaining, starting simple and building complexity.

## 1. Assign: Adding New Variables

**Mental Model:** Create new columns with calculations and transformations.

Let's start by calculating whether each shipment was late:

```
# Simple assignment - calculate if shipment was late
shipments_with_lateness = (
    shipments_df
    .assign(
        is_late=lambda df: df['actualShipDate'] > df['plannedShipDate'],
        days_late=lambda df: (df['actualShipDate'] - df['plannedShipDate']).dt.days
    )
)

print("Added lateness calculations:")
print(shipments_with_lateness[['shipID', 'plannedShipDate', 'actualShipDate', 'is_late', 'days_late']])
```

Added lateness calculations:

	shipID	plannedShipDate	actualShipDate	is_late	days_late
0	10001	2013-11-06	2013-10-04	False	-33
1	10002	2013-10-15	2013-10-04	False	-11
2	10003	2013-10-25	2013-10-07	False	-18
3	10004	2013-10-14	2013-10-08	False	-6
4	10005	2013-10-14	2013-10-08	False	-6

### Method Chaining Tip for New Python Users

**Why use `lambda df`?** When chaining methods, we need to reference the current state of the dataframe. The `lambda df` tells pandas “use the current dataframe in this calculation.” Without it, pandas would look for a variable called `df` that doesn’t exist.

**Alternative approach:** You could also write this as separate steps, but method chaining keeps related operations together and makes the code more readable.

### Discussion Questions: Assign Mental Model

**Question 1: Data Types and Date Handling** - What is the `dtype` of the `actualShipDate` series? How can you find out using code? - Why is it important that both `actualShipDate` and `plannedShipDate` have the same data type for comparison?

**Question 2: String vs Date Comparison** - Can you give an example where comparing two dates as strings would yield unintuitive results, e.g. what happens if you try to compare “04-11-2025” and “05-20-2024” as strings vs as dates?

### Question 3: Debug This Code

```
# This code has an error - can you spot it?
shipments_with_lateness = (
    shipments_df
    .assign(
        is_late=lambda df: df['actualShipDate'] > df['plannedShipDate'],
        days_late=lambda df: (df['actualShipDate'] - df['plannedShipDate']).dt.days,
        lateStatement="Darn Shipment is Late" if shipments_df['is_late'] else "Shipment is on time"
    )
)
```

What's wrong with the `lateStatement` assignment and how would you fix it?

### Briefly Give Answers to the Discussion Questions In This Section

#### Answer to Question 1:

- The dtype of the `actualShipDate` series is `datetime64[ns]`
- We can find out using various methods as follows:

```
# Method 1: Using .dtype attribute
print("actualShipDate dtype:", shipments_df['actualShipDate'].dtype)
print("=====")

# Method 2: Using .dtypes (for the entire dataframe)
print("All column dtypes:")
print(shipments_df.dtypes)
print("=====")

# Method 3: More detailed info with .info()
print("Detailed dataframe info:")
shipments_df.info()
print("=====")
```

```
actualShipDate dtype: datetime64[ns]
=====
All column dtypes:
shipID                int64
plannedShipDate       datetime64[ns]
actualShipDate        datetime64[ns]
partID                object
```

```

quantity                                int64
dtype: object
=====
Detailed dataframe info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   shipID                 4000 non-null   int64
1   plannedShipDate        4000 non-null   datetime64[ns]
2   actualShipDate          4000 non-null   datetime64[ns]
3   partID                 4000 non-null   object
4   quantity                4000 non-null   int64
dtypes: datetime64[ns](2), int64(2), object(1)
memory usage: 156.4+ KB
=====

```

- It is important that both `actualShipDate` and `plannedShipDate` have the same data type for comparison operations to work correctly. Pandas will compare the dates as strings if they are not the same type and comparisons would be alphabetical rather than chronological.

### Answer to Question 2:

When treated as strings, the comparison happens character by character.

- "04-11-2025" is compared to "05-20-2024".
- The first character 0 is the same.
- The second character 4 is compared to 5.
- Because 4 comes before 5 alphabetically, "04-11-2025" is considered "less than" "05-20-2024", which is wrong chronologically

```

from datetime import datetime
# The dates as strings
date_string_1 = "04-11-2025"
date_string_2 = "05-20-2024"

print(f"Comparing dates as strings:")
print(f"'{date_string_1}' > '{date_string_2}' is {date_string_1 > date_string_2}")
print(f"'{date_string_1}' < '{date_string_2}' is {date_string_1 < date_string_2}\n")

# The format of our dates is MM-DD-YYYY

```



```

date_format = "%m-%d-%Y"

# Convert the strings to datetime objects
date_object_1 = datetime.strptime(date_string_1, date_format)
date_object_2 = datetime.strptime(date_string_2, date_format)

print(f"Comparing dates as datetime objects:")
print(f"'{date_string_1}' > '{date_string_2}' is {date_object_1 > date_object_2}")
print(f"'{date_string_1}' < '{date_string_2}' is {date_object_1 < date_object_2}\n")

```

Comparing dates as strings:

```

'04-11-2025' > '05-20-2024' is False
'04-11-2025' < '05-20-2024' is True

```

Comparing dates as datetime objects:

```

'04-11-2025' > '05-20-2024' is True
'04-11-2025' < '05-20-2024' is False

```

### Answer to Question 3:

- lambda-df was missing to reference the dataframe
- lateStatement is a string, not a boolean so we need use vectorized operations to create the string like np.where()

```

#This is the corrected code
shipments_with_lateness = (
    shipments_df
    .assign(
        is_late=lambda df: df['actualShipDate'] > df['plannedShipDate'],
        days_late=lambda df: (df['actualShipDate'] - df['plannedShipDate']).dt.days,
        lateStatement=lambda df: np.where(df['is_late'], "Darn Shipment is Late", "Shipment :
    )
)

```

## 2. Subset: Querying Rows and Filtering Columns

**Mental Model:** Query rows based on conditions and filter to keep specific columns.

Let's query for only late shipments and filter to keep the columns we need:

```
# Query rows for late shipments and filter to keep specific columns
late_shipments = (
    shipments_with_lateness
    .query('is_late == True') # Query rows where is_late is True
    .filter(['shipID', 'partID', 'plannedShipDate', 'actualShipDate', 'days_late']) # Filter
)

print(f"Found {len(late_shipments)} late shipments out of {len(shipments_with_lateness)} total")
print("\nLate shipments sample:")
print(late_shipments.head())
```

Found 456 late shipments out of 4000 total

Late shipments sample:

	shipID	partID	plannedShipDate	actualShipDate	days_late
776	10192	part0164a70	2013-10-09	2013-10-14	5
777	10192	part9259836	2013-10-09	2013-10-14	5
778	10192	part4526c73	2013-10-09	2013-10-14	5
779	10192	partbb47e81	2013-10-09	2013-10-14	5
780	10192	part008482f	2013-10-09	2013-10-14	5

### Understanding the Methods

- **.query()**: Query rows based on conditions (like SQL WHERE clause)
- **.filter()**: Filter to keep specific columns by name
- **Alternative**: You could use **.loc[]** for more complex row querying, but **.query()** is often more readable

### Discussion Questions: Subset Mental Model

**Question 1: Query vs Boolean Indexing** - What's the difference between using **.query('is\_late == True')** and **[df['is\_late'] == True]**? - Which approach is more readable and why?

**Question 2: Additional Row Querying** - Can you show an example of using a variable like **late\_threshold** to query rows for shipments that are at least **late\_threshold** days late, e.g. what if you wanted to query rows for shipments that are at least 5 days late?

## Briefly Give Answers to the Discussion Questions In This Section

Answer to Question 1:

- In pandas, `df[df['is_late'] == True]` uses boolean indexing, while `df.query('is_late == True')` uses the `.query()` method. While they can produce the same result for simple cases, they operate differently and have trade-offs in terms of readability, performance, and flexibility.
- The `.query()` method is more readable and allows for more complex queries using SQL-like syntax.

## Answer to Question 2:

```
#Set the late threshold
late_threshold = 5

#Query the shipments that are at least 5 days late using the late_threshold variable
late_shipments = (
    shipments_with_lateness
    .query('days_late >= @late_threshold')
    .filter(['shipID', 'partID', 'plannedShipDate', 'actualShipDate', 'days_late'])
)
print(f"Found {len(late_shipments)} shipments that are at least {late_threshold} days late out of 4000 total")
```

Found 186 shipments that are at least 5 days late out of 4000 total

## 3. Drop: Removing Unwanted Data

**Mental Model:** Remove columns or rows you don't need.

Let's clean up our data by removing unnecessary columns:

```
# Create a cleaner dataset by dropping unnecessary columns
clean_shipments = (
    shipments_with_lateness
    .drop(columns=['quantity']) # Drop quantity column (not needed for our analysis)
    .dropna(subset=['plannedShipDate', 'actualShipDate']) # Remove rows with missing dates
)

print(f"Cleaned dataset: {len(clean_shipments)} rows, {len(clean_shipments.columns)} columns")
print("Remaining columns:", clean_shipments.columns.tolist())
```

Cleaned dataset: 4000 rows, 7 columns

Remaining columns: ['shipID', 'plannedShipDate', 'actualShipDate', 'partID', 'is\_late', 'days\_late']

## ! Discussion Questions: Drop Mental Model

**Question 1: Drop vs Filter Strategies** - What's the difference between `.drop(columns=['quantity'])` and `.filter()` with a list of columns you want to keep?

- When would you choose to drop columns vs filter to keep specific columns?

**Question 2: Handling Missing Data** - What happens if you use `.dropna()` without specifying subset? How is this different from `.dropna(subset=['plannedShipDate', 'actualShipDate'])`? - Why might you want to be selective about which columns to check for missing values?

### Briefly Give Answers to the Discussion Questions In This Section

#### Answer to Question 1:

- `.drop(columns=['quantity'])` removes that specific column from the dataframe. You would choose to drop columns if you don't need them for your analysis.
- `.filter()` explicitly specifies the columns you want to retain. You would create a new DataFrame by selecting only the columns present in your list.

#### Answer to Question 2:

- If you use `.dropna()` without specifying `subset`, it will remove all rows with any missing values in any column.
- This is different from `.dropna(subset=['plannedShipDate', 'actualShipDate'])` which will only remove rows with missing values in the `plannedShipDate` and `actualShipDate` columns.
- You might want to be selective about which columns to check for missing values if you have a lot of columns and you only want to check a few of them.

## 4. Sort: Arranging Data

**Mental Model:** Order data by values or indices.

Let's sort by lateness to see the worst offenders:

```
# Sort by days late (worst first)
sorted_by_lateness = (
    clean_shipments
    .sort_values('days_late', ascending=False) # Sort by days_late, highest first
    .reset_index(drop=True) # Reset index to be sequential
)
```

```
print("Shipments sorted by lateness (worst first):")
print(sorted_by_lateness[['shipID', 'partID', 'days_late', 'is_late']].head(10))
```

Shipments sorted by lateness (worst first):

	shipID	partID	days_late	is_late
0	10956	partb6208b5	21	True
1	10956	part04ef2f7	21	True
2	10956	part4875f85	21	True
3	10956	partb722d53	21	True
4	10956	partc979912	21	True
5	10956	parta27d449	21	True
6	10956	partc653823	21	True
7	10956	part82e69e9	21	True
8	10956	partf23fd1e	21	True
9	10956	part825873c	21	True

### ! Discussion Questions: Sort Mental Model

**Question 1: Sorting Strategies** - What's the difference between `ascending=False` and `ascending=True` in sorting? - How would you sort by multiple columns (e.g., first by `is_late`, then by `days_late`)?

**Question 2: Index Management** - Why do we use `.reset_index(drop=True)` after sorting? - What happens to the original index when you sort? Why might this be problematic?

## Briefly Give Answers to the Discussion Questions In This Section

### Answer to Question 1:

- `ascending=True`:
  - This sorts data in ascending order.
  - For numerical values, it means from smallest to largest (e.g., 1, 2, 3).
  - For alphabetical values, it means from A to Z (e.g., Apple, Banana, Cherry).
  - For dates, it means from earliest to latest (e.g., 2023-01-01, 2023-01-02).
- `ascending=False`:
  - This sorts data in descending order.
  - For numerical values, it means from largest to smallest (e.g., 3, 2, 1).
  - For alphabetical values, it means from Z to A (e.g., Cherry, Banana, Apple).
  - For dates, it means from latest to earliest (e.g., 2023-01-02, 2023-01-01).

- To sort by multiple columns, such as first by `is_late` and then by `days_late`, you typically provide a list of column names to the sorting function or method. The order of columns in this list determines the priority of sorting. The first column in the list is the primary sort key, and subsequent columns are used to break ties within the primary sort.

```
# Sort first by is_late (False first, then True), then by days_late
sorted_shipments = (
    clean_shipments
    .sort_values(
        by=['is_late', 'days_late'],
        ascending=[True, False] # is_late ascending (False first), days_late descending
    )
    .reset_index(drop=True)
)

print("Shipments sorted by first is_late and then by days_late :")
print(sorted_shipments[['shipID', 'partID', 'days_late', 'is_late']].head(50))
```

Shipments sorted by first is\_late and then by days\_late :

	shipID	partID	days_late	is_late
0	10193	part71b107e	0	False
1	10193	part8acdaf4	0	False
2	10193	part3abf99d	0	False
3	10193	part6ee349a	0	False
4	10193	part7259bf2	0	False
5	10193	partfe2de18	0	False
6	10193	part82548d0	0	False
7	10193	partb60c5be	0	False
8	10193	partf194391	0	False
9	10193	partb8509b1	0	False
10	10193	parte04a6ce	0	False
11	10193	partfcb753c	0	False
12	10193	part0e8d20e	0	False
13	10193	part2302c27	0	False
14	10193	partd578ab3	0	False
15	10193	part15541ff	0	False
16	10193	part6527169	0	False
17	10193	part9146d41	0	False
18	10193	part57669fb	0	False
19	10193	part21bade2	0	False
20	10193	partba76e17	0	False

21	10193	partef9a92a	0	False
22	10193	parta2cc65f	0	False
23	10193	part481932d	0	False
24	10193	partcbde9e0	0	False
25	10193	part472d67c	0	False
26	10193	part7547885	0	False
27	10193	parta442b41	0	False
28	10193	partf557daa	0	False
29	10207	part0e9dcf0	0	False
30	10207	part283dfaa	0	False
31	10207	partcfdcff7	0	False
32	10207	part5f49e4d	0	False
33	10207	part0c5bf9a	0	False
34	10207	part5107513	0	False
35	10207	part05257e	0	False
36	10207	part07bee5e	0	False
37	10207	part08d5f32	0	False
38	10207	parte310d1	0	False
39	10207	part2e3c85b	0	False
40	10207	part8fb9837	0	False
41	10207	part8100a77	0	False
42	10207	part74dc9f	0	False
43	10207	part198d6c2	0	False
44	10207	part1d3bfad	0	False
45	10207	partfd372fb	0	False
46	10207	part1d8d910	0	False
47	10207	partc1987d8	0	False
48	10207	part07671a5	0	False
49	10207	parteea7762	0	False

First sort by is\_late:

- False (on-time) shipments come first
- True (late) shipments come second

Then sort by days\_late within each group:

- For on-time shipments: most negative (earliest) first
- For late shipments: most positive (latest) first

**Answer to Question 2:**

- Using `.reset_index(drop=True)` after sorting a Pandas DataFrame serves to re-establish a default, sequential integer index while discarding the old, potentially non-sequential or non-numeric index.
- Maintaining a Clean, Sequential Index: When a DataFrame is sorted, the original index labels remain associated with their respective rows, even though the order of the rows changes. This can lead to a non-sequential or jumbled index. `reset_index()` replaces this with a new, clean index starting from 0.
- Facilitating Positional Indexing: A sequential integer index makes it easier to access rows by their position using `iloc`, as the index directly corresponds to the row number.
- Preventing Misleading Index Values: If the original index held no inherent meaning beyond identifying the row's original position, retaining it after sorting can be misleading, as the index values no longer reflect the current order of the data.
- `drop=True`'s Role: The `drop=True` argument is crucial here because it prevents the old index from being added as a new column in the DataFrame. If `drop=True` were omitted, the old index would be converted into a regular column, which is often not the desired behavior after sorting.

## 5. Aggregate: Summarizing Data

**Mental Model:** Calculate summary statistics across groups or the entire dataset.

Let's calculate overall service level metrics:

```
# Calculate overall service level metrics
service_metrics = (
    clean_shipments
    .agg({
        'is_late': ['count', 'sum', 'mean'], # Count total, count late, calculate percentage
        'days_late': ['mean', 'max'] # Average and maximum days late
    })
    .round(3)
)

print("Overall Service Level Metrics:")
print(service_metrics)

# Calculate percentage on-time directly from the data
on_time_rate = (1 - clean_shipments['is_late'].mean()) * 100
print(f"\nOn-time delivery rate: {on_time_rate:.1f}%")
```

```
Overall Service Level Metrics:
      is_late  days_late
```



count	4000.000	NaN
sum	456.000	NaN
mean	0.114	-0.974
max	NaN	21.000

On-time delivery rate: 88.6%

### ! Discussion Questions: Aggregate Mental Model

**Question 1: Boolean Aggregation** - Why does `sum()` work on boolean values? What does it count?

## Briefly Give Answers to the Discussion Questions In This Section

### Answer to Q1:

- In Pandas, the `sum()` works on boolean values because `bool` is a subclass of `int` in Python. This means that `True` is treated as 1 and `False` is treated as 0 when numerical operations are performed.
- When we call `sum()` on a Series or DataFrame column containing boolean values, it effectively counts the number of `True` values present. Each `True` value contributes 1 to the sum, while each `False` value contributes 0.

## 6. Merge: Combining Information

**Mental Model:** Join data from multiple sources to create richer datasets.

Now let's analyze service levels by product category. First, we need to merge our data:

```
# Merge shipment data with product line data
shipments_with_category = (
    clean_shipments
    .merge(product_line_df, on='partID', how='left') # Left join to keep all shipments
    .assign(
        category_late=lambda df: df['is_late'] & df['prodCategory'].notna() # Only count as
    )
)

print("\nProduct categories available:")
print(shipments_with_category['prodCategory'].value_counts())
```

Product categories available:

```
prodCategory
Marketables    1850
Machines        846
SpareParts     767
Liquids         537
Name: count, dtype: int64
```

**!** Discussion Questions: Merge Mental Model

**Question 1: Join Types and Data Loss** - Why does your professor think we should use `how='left'` in most cases? - How can you check if any shipments were lost during the merge?

**Question 2: Key Column Matching** - What happens if there are duplicate `partID` values in the `product_line_df`?

### Briefly Give Answers to the Discussion Questions In This Section

#### Answer to Q1:

- Using `how='left'` in `.merge()` is often preferred because it guarantees that all rows from the “left” DataFrame are preserved in the merged result. `how='left'` is a common and safe default when the goal is to add information to an existing, primary dataset without losing any of its original records. This is crucial in many data analysis scenarios for several reasons:
  - **Data Integrity:** It ensures that no data from your primary (left) DataFrame is accidentally dropped due to a lack of matching keys in the “right” DataFrame. This helps maintain the integrity of your original dataset.
  - **Adding Contextual Information:** When you want to enrich an existing dataset with additional information from another source, a left merge allows you to add those details without losing any of your original observations. If a match isn’t found in the right DataFrame, the corresponding columns will simply be filled with NaN values, indicating the absence of that specific information for those rows.
  - **Predictability and Control:** Left merges offer more predictable behavior regarding the number of rows in the output. The resulting DataFrame will always contain at least as many rows as the left DataFrame.
  - **Debugging and Traceability:** If you need to investigate why certain rows lack information after a merge, the NaN values resulting from a left merge clearly indicate which rows did not find a match in the right DataFrame. This makes debugging and understanding your data easier.

To check if any shipment were lost during the merge, we can verify the using various ways such as :

- Compare before / after row counts.- Check for any rows with missing values.
- Check for any rows having duplicate values.

Here is an example code for Comparing row counts:

```
# Check row counts before and after merge
print(f"Original shipments: {len(clean_shipments)}")
print(f"After merge: {len(shipments_with_category)}")
print(f"Difference: {len(shipments_with_category) - len(clean_shipments)}")

# Verify no data loss
if len(shipments_with_category) == len(clean_shipments):
    print("No shipments lost during merge")
else:
    print("Data loss detected!")
```

```
Original shipments: 4000
After merge: 4000
Difference: 0
No shipments lost during merge
```

### Answer to Q2:

If there are duplicate `partID` values in the `product_line_df`, we can run into data integrity issues. In his scenario, merge will create multiple rows for each shipment that matches those duplicate `partIDs`.

Here is an example:

- `product_line_df` has :

partID	productLine	productCategory
part123	Line1	Machines
part123	Line2	Liquids #Duplicate partID
- While merging with `shipments_df`:

Each shipment with `partID` “part123” will get BOTH rows. As a Result, 1 shipment becomes 2 rows in the merged data.

## 7. Split-Apply-Combine: Group Analysis

**Mental Model:** Group data and apply functions within each group.

Now let's analyze service levels by category:

```
# Analyze service levels by product category
service_by_category = (
    shipments_with_category
    .groupby('prodCategory') # Split by product category
    .agg({
        'is_late': ['any', 'count', 'sum', 'mean'], # Count, late count, percentage late
        'days_late': ['mean', 'max'] # Average and max days late
    })
    .round(3)
)

print("Service Level by Product Category:")
print(service_by_category)
```

Service Level by Product Category:

	is_late			days_late		
prodCategory	any	count	sum	mean	mean	max
Liquids	True	537	22	0.041	-0.950	19
Machines	True	846	152	0.180	-1.336	21
Marketable	True	1850	145	0.078	-0.804	21
SpareParts	True	767	137	0.179	-1.003	21

### ! Discussion Questions: Split-Apply-Combine Mental Model

**Question 1: GroupBy Mechanics** - What does `.groupby('prodCategory')` actually do? How does it “split” the data? - Why do we need to use `.agg()` after grouping? What happens if you don't?

**Question 2: Multi-Level Grouping** - Explore grouping by `['shipID', 'prodCategory']`? What question does this answer versus grouping by `'prodCategory'` alone? (HINT: There may be many rows with identical `shipID`'s due to a particular order having multiple `partID`'s.)

## Briefly Give Answers to the Discussion Questions In This Section

### Answer to Q1:

- In pandas, groupby combines “split-apply-combine” process into one with the help of `.agg()`.
- Just `.groupby('prodCategory')` will create a GroupBy object that contains instructions for how to split the data based on a ‘productCategory’. The groupby operation then creates separate data chunks for each unique value in grouping column `productCategory`.
- You need `.agg()` after grouping to perform specific calculations (like sum, mean, max) on each group, reducing the data to a summarized, single row per group. It allows you to apply one or more aggregation functions to a DataFrame, making it much easier to summarize and analyze data. Without an aggregation function, the grouping operation itself would not produce a useful, summarized result.
- If you group a DataFrame but do not specify an aggregation, you are left with a “GroupBy object,” which is not the final result you typically want. This object is an intermediate step in the “split-apply-combine” process, and you must apply an aggregation function (like `.agg()`, `.sum()`, or `.mean()`) to the “apply” and “combine” those results into a meaningful output.

### Answer to Q2:

Grouping by 'prodCategory' alone:

- Question: "What is the overall service level for each product category?"
- Answer: Service metrics aggregated across ALL shipments for each category
- How many total shipments in each category?
- What percentage of shipments are late in each category?
- Average and maximum lateness by category

Grouping by ['shipID', 'prodCategory']:

- Question: "What is the service level for each shipment within each product category?"
- Answer: Service metrics for each individual shipment within each category
- Each row represents one shipment's performance in one category
- Shows which specific shipments had issues in which categories

## Answering A Business Question

**Mental Model:** Combine multiple data manipulation techniques to answer complex business questions.

Let's create a comprehensive analysis by combining shipment-level data with category information:

```
# Create a comprehensive analysis dataset
comprehensive_analysis = (
    shipments_with_category
    .groupby(['shipID', 'prodCategory']) # Group by shipment and category
    .agg({
        'is_late': 'any', # True if any item in this shipment/category is late
        'days_late': 'max' # Maximum days late for this shipment/category
    })
    .reset_index()
    .assign(
        has_multiple_categories=lambda df: df.groupby('shipID')['prodCategory'].transform('n
    )
)

print("Comprehensive analysis - shipments with multiple categories:")
multi_category_shipments = comprehensive_analysis[comprehensive_analysis['has_multiple_categ
print(f"Shipments with multiple categories: {multi_category_shipments['shipID'].nunique()}")
print(f"Total unique shipments: {comprehensive_analysis['shipID'].nunique()}")
print(f"Percentage with multiple categories: {multi_category_shipments['shipID'].nunique() /
```

Comprehensive analysis - shipments with multiple categories:  
Shipments with multiple categories: 232  
Total unique shipments: 997  
Percentage with multiple categories: 23.3%

### ! Discussion Questions: Answering A Business Question

**Question 1: Business Question Analysis** - What business question does this comprehensive analysis answer? - How does grouping by ['shipID', 'prodCategory'] differ from grouping by just 'prodCategory'? - What insights can ZappTech's management gain from knowing the percentage of multi-category shipments?

### Briefly Give Answers to the Discussion Questions In This Section

#### Answer to Q1:

The Business Question This Comprehensive Analysis Answers Primary Question:

- “How often do our customers place orders that span multiple product categories, and what are the operational implications of these cross-category orders?”

- Grouping by ['shipID', 'prodCategory'] creates one row per shipmentID-prodCategory combination as compared to just prodCategory which would have created one row each per category only.
- It then identifies shipments that contain products from more than one category.
- Finally, shows what amount of all shipments are multi-category as percentage.

Business insights ZappTech's management can know from percentage of multi-category shipments:

- Cross-Category Shopping behavior: 232 out of 997 (23.3%) contain multiple productCategories. This tells about customer ordering patterns.

Business Question: "Do our customers typically order from one category at a time, or do they frequently mix categories in single orders?"

- Operational Complexity

Business Question: "What percentage of our shipments require coordination across multiple product lines, and how does this affect our fulfillment process?"

- Service Level Impact

Business Question: "Are multi-category orders more likely to experience delays due to the complexity of coordinating different product lines?"

## Student Analysis Section: Mastering Data Manipulation

**Your Task:** Demonstrate your mastery of the seven mental models through comprehensive discussion and analysis. The bulk of your grade comes from thoughtfully answering the discussion questions for each mental model. See below for more details.

### Core Challenge: Discussion Questions Analysis

**For each mental model, provide:** - Clear, concise answers to all discussion questions - Code examples where appropriate to support your explanations

#### ! Discussion Questions Requirements

**Complete all discussion question sections:** 1. **Assign Mental Model:** Data types, date handling, and debugging 2. **Subset Mental Model:** Filtering strategies and complex queries 3. **Drop Mental Model:** Data cleaning and quality management 4. **Sort Mental Model:** Data organization and business logic 5. **Aggregate Mental Model:** Summary statistics and business metrics 6. **Merge Mental Model:** Data integration and quality control 7. **Split-Apply-Combine Mental Model:** Group

analysis and advanced operations 8. **Answering A Business Question:** Combining multiple data manipulation techniques to answer a business question

## Professional Visualizations (For 100% Grade)

**Your Task:** Create a professional visualization that supports your analysis and demonstrates your understanding of the data.

**Create visualizations showing:** - Service level (on-time percentage) by product category

**Your visualizations should:** - Use clear labels and professional formatting - Support the insights from your discussion questions - Be appropriate for a business audience - Do not echo the code that creates the visualizations

## Challenge Requirements

**Your Primary Task:** Answer all discussion questions for the seven mental models with thoughtful, well-reasoned responses that demonstrate your understanding of data manipulation concepts.

**Key Requirements:** - Complete discussion questions for each mental model - Demonstrate clear understanding of pandas concepts and data manipulation techniques - Write clear, business-focused analysis that explains your findings

## Getting Started: Repository Setup

### ! Getting Started

**Step 1:** Fork and clone this challenge repository - Go to the course repository and find the “dataManipulationChallenge” folder - Fork it to your GitHub account, or clone it directly - Open the cloned repository in Cursor

**Step 2:** Set up your Python environment - Follow the Python setup instructions above (use your existing venv from Tech Setup Challenge Part 2) - Make sure your virtual environment is activated and the Python interpreter is set

**Step 3:** You’re ready to start! The data loading code is already provided in this file.

**Note:** This challenge uses the same `index.qmd` file you’re reading right now - you’ll edit it to complete your analysis.



## Getting Started Tips

### Method Chaining Philosophy

“Each operation should build naturally on the previous one”

*Think of method chaining like building with LEGO blocks - each piece connects to the next, creating something more complex and useful than the individual pieces.*

### Important: Save Your Work Frequently!

**Before you start:** Make sure to commit your work often using the Source Control panel in Cursor (Ctrl+Shift+G or Cmd+Shift+G). This prevents the AI from overwriting your progress and ensures you don't lose your work.

**Commit after each major step:**

- After completing each mental model section
- After adding your visualizations
- After completing your advanced method chain
- Before asking the AI for help with new code

**How to commit:**

1. Open Source Control panel (Ctrl+Shift+G)
2. Stage your changes (+ button)
3. Write a descriptive commit message
4. Click the checkmark to commit

*Remember: Frequent commits are your safety net!*

## Grading Rubric

**75% Grade:** Complete discussion questions for at least 5 of the 7 mental models with clear, thoughtful responses.

**85% Grade:** Complete discussion questions for all 7 mental models with comprehensive, well-reasoned responses.

**95% Grade:** Complete all discussion questions plus the “Answering A Business Question” section.

**100% Grade:** Complete all discussion questions plus create a professional visualization showing service level by product category.

## Submission Checklist

### Minimum Requirements (Required for Any Points):

- ☐ Created repository named “dataManipulationChallenge” in your GitHub account
- ☐ Cloned repository locally using Cursor (or VS Code)
- ☐ Completed discussion questions for at least 5 of the 7 mental models
- ☐ Document rendered to HTML successfully
- ☐ HTML files uploaded to your repository
- ☐ GitHub Pages enabled and working
- ☐ Site accessible at [https://\[your-username\].github.io/dataManipulationChallenge/](https://[your-username].github.io/dataManipulationChallenge/)

### 75% Grade Requirements:

- ☐ Complete discussion questions for at least 5 of the 7 mental models
- ☐ Clear, thoughtful responses that demonstrate understanding
- ☐ Code examples where appropriate to support explanations

### 85% Grade Requirements:

- ☐ Complete discussion questions for all 7 mental models
- ☐ Comprehensive, well-reasoned responses showing deep understanding
- ☐ Business context for why concepts matter
- ☐ Examples of real-world applications

### 95% Grade Requirements:

- ☐ Complete discussion questions for all 7 mental models
- ☐ Complete the “Answering A Business Question” discussion questions
- ☐ Comprehensive, well-reasoned responses showing deep understanding
- ☐ Business context for why concepts matter

### 100% Grade Requirements:

- ☐ All discussion questions completed with professional quality
- ☐ Professional visualization showing service level by product category
- ☐ Professional presentation style appropriate for business audience
- ☐ Clear, engaging narrative that tells a compelling story
- ☐ Practical insights that would help ZappTech’s management

### Report Quality (Critical for Higher Grades):

- ☐ Professional writing style (no AI-generated fluff)
- ☐ Concise analysis that gets to the point